

# LALE: Consistent Automated Machine Learning

---

Guillaume Baudart  
Martin Hirzel  
Kiran Kate  
Pari Ram  
Avi Shinnar



AutoML Global Summit  
Sep 3, 2020

<https://github.com/ibm/lale>

---

**IBM Research**



# Spectrum of Automation

|                             |                          | What is being configured?                            |                              |                   |
|-----------------------------|--------------------------|------------------------------------------------------|------------------------------|-------------------|
|                             |                          | Operator choice                                      | Hyperparameter configuration | Pipeline topology |
| How is it being configured? | Manual                   | scikit-learn                                         |                              |                   |
|                             | Automated (with control) | sklearn.model_selection.GridSearchCV, SMAC, hyperopt |                              | AlphaD3M          |
|                             | Automated (black-box)    | auto-sklearn, hyperopt-sklearn                       |                              | TPOT              |

# Existing AutoML tools ...

- ... automatically find good configurations 😊
- ... are open-source 😊
- ... are scikit-learn-based 😊
- ... support only part of the automation spectrum
- ... use inconsistent programming models

# Programming Model Requirements

- Support entire automation spectrum in a consistent way
- Be consistent across tools (hyperopt, GridSearchCV, SMAC)
  - ➔ Search space generation
- Extend established abstractions (scikit-learn, JSON schema)

# Operator Choice

## Manual

```
pipeline = J48()  
trained = pipeline.fit(X, y)
```

## Automated (aka. algorithm selection)

```
pipeline = J48() | LR()  
trained = pipeline.auto_configure(X, y, optimizer=GridSearchCV)
```

# Hyperparameter Configuration

## Manual

```
pipeline = J48(R=False, C=0.3)
trained = pipeline.fit(X, y)
```

## Automated (aka. hyperparameter tuning)

```
pipeline = J48
trained = pipeline.auto_configure(X, y, optimizer=SMAC)
```

```
J48: { allOf: [
  { type: object,
    properties: {
      R: { description: "Use reduced error pruning", type: boolean },
      C: { description: "Pruning confidence threshold",
        type: number, minimum: 0.0, maximum: 1.0,
        maximumForOptimizer: 0.5, distribution: uniform }},
    { description: "Setting confidence makes no sense for R",
      anyOf: [
        { not: { type: object, properties: { R: { enum: [true] } } } },
        { type: object, properties: { C: { enum: [0.25] } } } } ]
    }
  ]
}
```

# Pipeline Composition

## Manual

```
prep_num = Project(columns={'type':'number'}) >> PCA()  
prep_str = Project(columns={'type':'string'}) >> OneHotEncoder()  
pipeline = (prep_num & prep_str) >> ConcatFeatures >> LR()  
trained = pipeline.fit(X, y)
```

## Automated (aka. topology search)

```
g.start = g.prep >> (J48 | LR)  
g.prep = NoOp | (g.prep >> g.prep1)  
g.prep1 = StandardScaler | Normalizer | PolynomialFeatures | PCA  
trained = g.unfold(3).auto_configure(X, y, optimizer=Hyperopt)
```

# Higher-Order Operators

## Manual

```
tree = DecisionTreeClassifier(max_depth=1)
clf = AdaBoostClassifier(base_estimator=tree, n_estimators=10)
pipeline = StandardScaler() >> clf
trained = pipeline.fit(X, y)
```

## Automated (aka. topology search)

```
clf = AdaBoostClassifier(base_estimator=DecisionTreeClassifier)
pipeline = (StandardScaler | PCA | NoOp) >> clf
trained = pipeline.auto_configure(X, y, optimizer=Hyperopt)
```



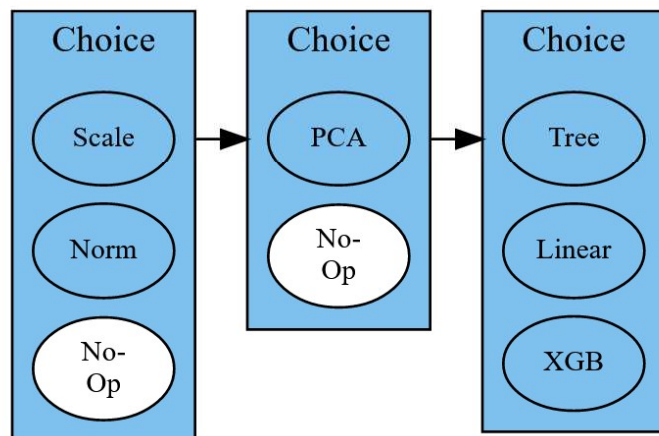
# Demonstration

```
In [2]: ▶ 1 from sklearn.preprocessing import StandardScaler as Scale
          2 from sklearn.preprocessing import Normalizer as Norm
          3 from lale.lib.lale import NoOp
          4 from sklearn.decomposition import PCA
          5 from sklearn.tree import DecisionTreeRegressor as Tree
          6 from sklearn.linear_model import LinearRegression as Linear
          7 from xgboost import XGBRegressor as XGB
          8 lale.wrap_imported_operators()
```

# Demonstration

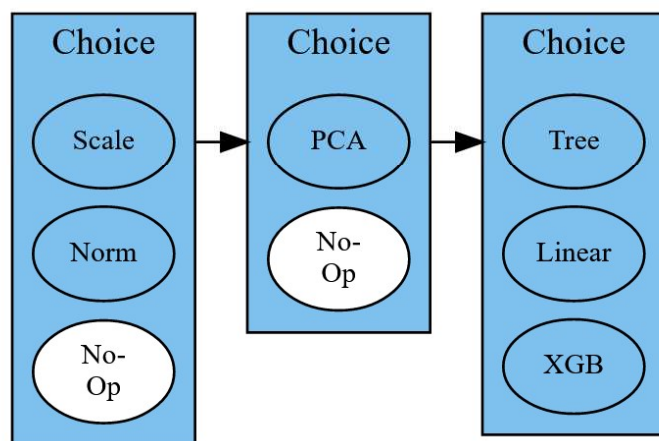
```
In [2]: ▶ 1 from sklearn.preprocessing import StandardScaler as Scale
          2 from sklearn.preprocessing import Normalizer as Norm
          3 from lale.lib.lale import NoOp
          4 from sklearn.decomposition import PCA
          5 from sklearn.tree import DecisionTreeRegressor as Tree
          6 from sklearn.linear_model import LinearRegression as Linear
          7 from xgboost import XGBRegressor as XGB
          8 lale.wrap_imported_operators()
```

```
In [3]: ▶ 1 planned_pipeline = (Scale | Norm | NoOp) >> (PCA | NoOp) >> (Tree | Linear | XGB)
          2 planned_pipeline.visualize()
```



# Demonstration

```
In [3]: ► 1 planned_pipeline = (Scale | Norm | NoOp) >> (PCA | NoOp) >> (Tree | Linear | XGB)
2 planned_pipeline.visualize()
```



```
In [4]: ► 1 from lale.lib.lale import Hyperopt
2 import sklearn.metrics
3 r2 = sklearn.metrics.make_scorer(sklearn.metrics.r2_score)
4 trained_pipeline = planned_pipeline.auto_configure(
5     train_X, train_y, optimizer=Hyperopt,
6     scoring=r2, max_opt_time=10*60, max_eval_time=60, cv=3)
```

```
100%|██████████| 50/50 [10:02<00:00, 12.05s/trial, best loss: -0.8107292214446913]
```


# Demonstration

```
In [4]: 1 from lale.lib.lale import Hyperopt
2 import sklearn.metrics
3 r2 = sklearn.metrics.make_scorer(sklearn.metrics.r2_score)
4 trained_pipeline = planned_pipeline.auto_configure(
5     train_X, train_y, optimizer=Hyperopt,
6     scoring=r2, max_opt_time=10*60, max_eval_time=60, cv=3)

100%|██████████| 50/50 [10:02<00:00, 12.05s/trial, best loss: -0.8107292214446913]
```

```
In [5]: 1 print(f'R2 score: {r2(trained_pipeline, test_X, test_y):.2f}')
2 trained_pipeline.visualize()
```

R2 score: 0.83



# Demonstration

```
In [5]: ▶ 1 print(f'R2 score: {r2(trained_pipeline, test_X, test_y):.2f}')
          2 trained_pipeline.visualize()
```

R2 score: 0.83

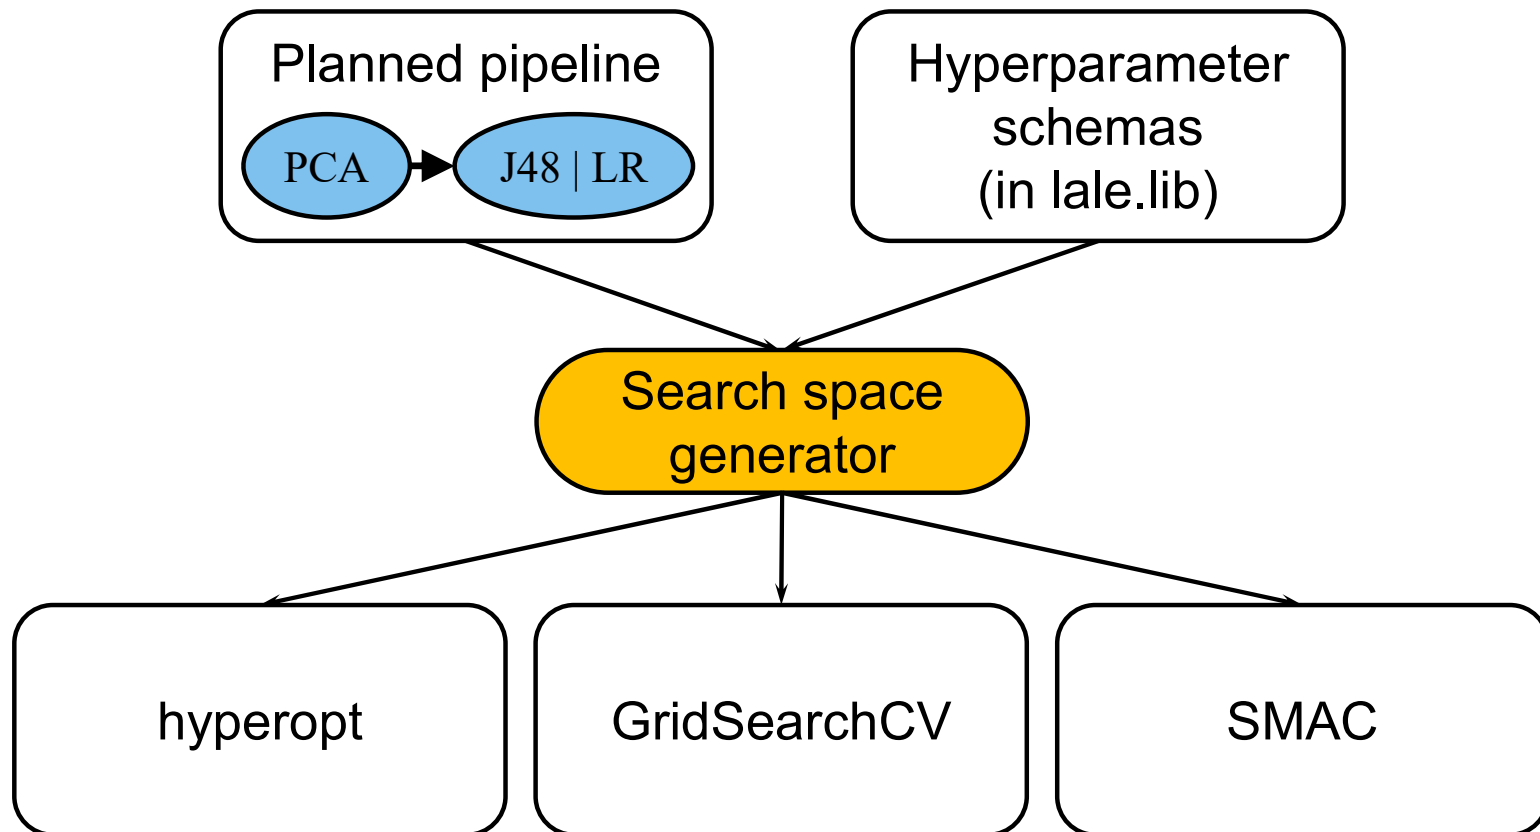


```
In [6]: ▶ 1 trained_pipeline.pretty_print(ipython_display=True)
```

```
from lale.lib.sklearn.standard_scaler import Scale
from lale.lib.lale import NoOp
from lale.lib.xgboost.xgb_regressor import XGB
import lale
lale.wrap_imported_operators()
```

```
xgb = XGB(colsample_bylevel=0.8539909881709349, colsample_bytree=0.9733482608060157, learning_rate=0.33190877144882
86, max_depth=10, min_child_weight=14, n_estimators=819, reg_alpha=0.07519396924044132, reg_lambda=0.08746289842357
71, subsample=0.9433039199868445)
pipeline = Scale() >> NoOp() >> xgb
```

# Consistency Across Existing Tools



# Normalize

## Hyperparameter schemas (in lale.lib)

*PCA* : dict{*N*: (0..1)  $\vee$  [*mle*]}

*J48* : dict{*R*: [*true*, *false*], *C*: (0..0.5)}  $\wedge$   
(dict{*R*: [*true*]}  $\Rightarrow$  dict{*C*: [0.25]})

*LR* : dict{*S*: [*linear*, *sag*, *lbfgs*], *P*: [*l1*, *l2*]}  $\wedge$   
(dict{*S*: [*sag*, *lbfgs*]}  $\Rightarrow$  dict{*P*: [*l2*]})



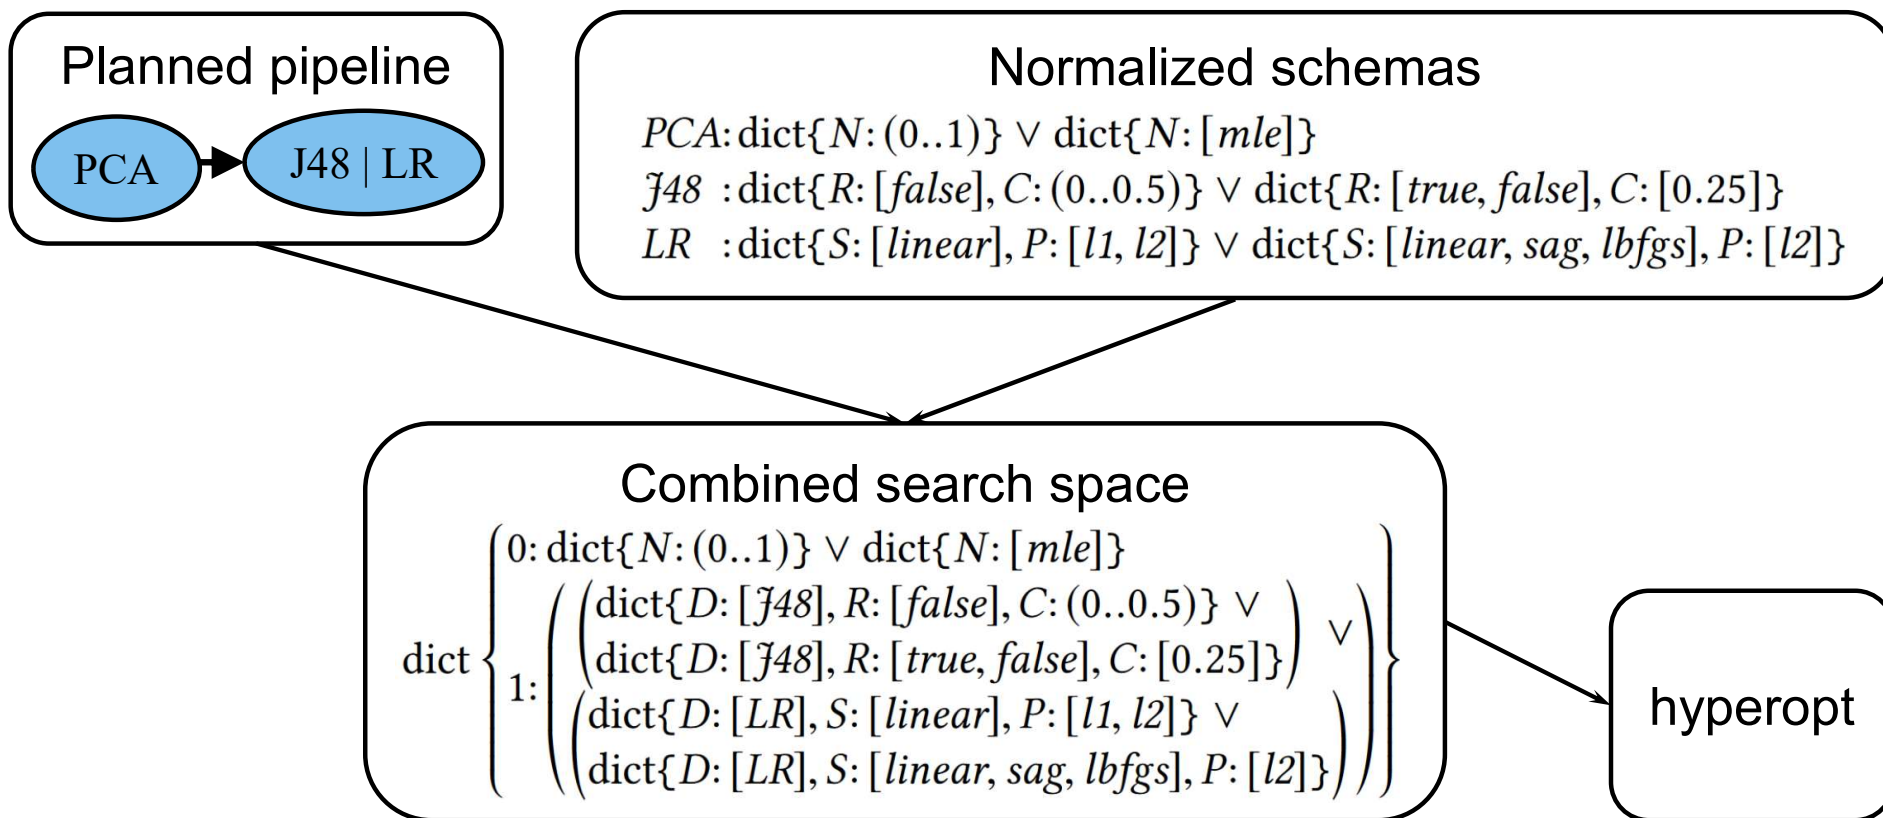
## Normalized schemas

*PCA*: dict{*N*: (0..1)}  $\vee$  dict{*N*: [*mle*]}

*J48* : dict{*R*: [*false*], *C*: (0..0.5)}  $\vee$  dict{*R*: [*true*, *false*], *C*: [0.25]}


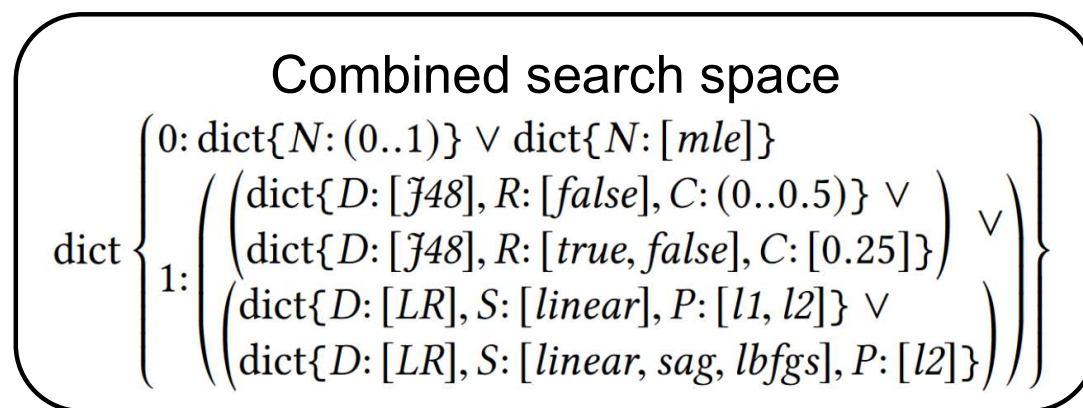
*LR* : dict{*S*: [*linear*], *P*: [*l1*, *l2*]}  $\vee$  dict{*S*: [*linear*, *sag*, *lbfgs*], *P*: [*l2*]}

# Combine





# Flatten



# Discretize

## Flat search space

$\text{dict}\{N: (0..1), D: [\text{J48}], R: [\text{false}], C: (0..0.5)\}$   
 $\vee \text{dict}\{N: (0..1), D: [\text{J48}], R: [\text{true}, \text{false}], C: [0.25]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{J48}], R: [\text{false}], C: (0..0.5)\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{J48}], R: [\text{true}, \text{false}], C: [0.25]\}$   
 $\vee \text{dict}\{N: (0..1), D: [\text{LR}], S: [\text{linear}], P: [l1, l2]\}$   
 $\vee \text{dict}\{N: (0..1), D: [\text{LR}], S: [\text{linear}, \text{sag}, \text{lbfgs}], P: [l2]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{LR}], S: [\text{linear}], P: [l1, l2]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{LR}], S: [\text{linear}, \text{sag}, \text{lbfgs}], P: [l2]\}$

## Discretized search space

$\text{dict}\{N: [0.50, 0.01], D: [\text{J48}], R: [\text{false}], C: [0.25, 0.01]\}$   
 $\vee \text{dict}\{N: [0.50, 0.01], D: [\text{J48}], R: [\text{true}, \text{false}], C: [0.25]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{J48}], R: [\text{false}], C: [0.25, 0.01]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{J48}], R: [\text{true}, \text{false}], C: [0.25]\}$   
 $\vee \text{dict}\{N: [0.50, 0.01], D: [\text{LR}], S: [\text{linear}], P: [l1, l2]\}$   
 $\vee \text{dict}\{N: [0.50, 0.01], D: [\text{LR}], S: [\text{linear}, \text{sag}, \text{lbfgs}], P: [l2]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{LR}], S: [\text{linear}], P: [l1, l2]\}$   
 $\vee \text{dict}\{N: [\text{mle}], D: [\text{LR}], S: [\text{linear}, \text{sag}, \text{lbfgs}], P: [l2]\}$

Grid-  
Search-  
CV

# 179 Operators with Schemas

| Package                  | Count | Description                                                      |
|--------------------------|-------|------------------------------------------------------------------|
| lale.lib.sklearn         | 46    | Hand-curated scikit-learn operators (subset of lale.lib.autogen) |
| lale.lib.autogen         | 115   | Auto-extracted scikit-learn operators                            |
| lale.lib.aif360          | 4     | Fairness mitigator                                               |
| lale.lib.autoai_libs     | 21    | Data cleansing and feature engineering                           |
| github.com/lale/lale-gpl | 2     | J48 from Weka and ARulesCBA from R                               |
| lale.lib.imblearn        | 12    | Class imbalance handling                                         |
| lale.lib.lale            | 15    | Optimizers and utility operators                                 |
| lale.lib.lightgbm        | 2     | Gradient-boosted random forests                                  |
| lale.lib.pai4sk          | 2     | Snap ML high-speed training                                      |
| lale.lib.pytorch         | 2     | BERT and ResNet                                                  |
| lale.lib.spacy           | 1     | Glove                                                            |
| lale.lib.tensorflow      | 1     | Universal sentence encoder                                       |
| lale.lib.xgboost         | 2     | Gradient-boosted random forests                                  |

# Results

- See paper  
<https://arxiv.org/abs/2007.01977>
- Comparison against auto-sklearn
  - 4 Lale pipelines, including grammars and higher-order operators
  - Competitive results on 15 OpenML datasets
- Case studies with other modalities
  - Text, images, time-series
- Effects of side constraints on convergence
  - Pruning the search space beats catching exceptions from trials

| DATASET       | 100 * ( <i>accuracy</i> /AUTOSKL - 1) |           |           |          |
|---------------|---------------------------------------|-----------|-----------|----------|
|               | LALE-PIPE                             | LALE-TPOT | LALE-AD3M | LALE-ADB |
| australian    | 0.41                                  | 0.93      | 2.06      | 1.13     |
| blood         | -2.08                                 | -0.52     | -4.05     | -1.04    |
| breast-cancer | -2.59                                 | -2.31     | -4.90     | -2.88    |
| car           | -1.13                                 | -0.25     | -6.70     | -1.09    |
| credit-g      | -2.29                                 | -3.24     | -2.37     | -0.71    |
| diabetes      | 0.61                                  | -0.82     | 1.12      | -1.33    |
| hill-valley   | -0.20                                 | 0.55      | -2.66     | 0.55     |
| jungle-chess  | 2.54                                  | 0.96      | -15.80    | 1.53     |
| kc1           | -0.38                                 | -0.38     | -0.21     | -0.58    |
| kr-vs-kp      | -0.36                                 | -0.27     | -2.87     | -0.19    |
| mfeat-factors | -1.14                                 | -1.54     | -1.17     | -1.20    |
| phoneme       | -1.39                                 | -0.83     | -15.20    | -0.22    |
| shuttle       | 14.51                                 | 14.50     | 14.45     | 14.56    |
| spectf        | -0.78                                 | 0.59      | -4.90     | 0.59     |
| sylvine       | -0.45                                 | -1.07     | -4.31     | -0.29    |

# Spectrum of Automation, Revisited

|                             |                          | What is being configured?                           |                              |                   |
|-----------------------------|--------------------------|-----------------------------------------------------|------------------------------|-------------------|
|                             |                          | Operator choice                                     | Hyperparameter configuration | Pipeline topology |
| How is it being configured? | Manual                   | scikit-learn                                        |                              |                   |
|                             | Automated (with control) | sklearn.model_selection.GridSearchCV, MAC, hyperopt |                              |                   |
|                             | Automated (black-box)    | auto-sklearn, hyperopt-sklearn                      |                              | TPOT              |

**L A L E**

# Conclusion

- Lale is open-source
  - <https://github.com/ibm/lale>
  - Try it
  - Use it
  - Contribute
- Lale is used by IBM AutoAI
  - <https://www.ibm.com/cloud/watson-studio/autoai>
  - Find and deploy top-performing models in minutes