

# Fairness for Scikit-Learn Pipelines with Lale

Martin Hirzel, IBM Research

PyData NYC November 2022

# Motivation

- Goal: fair machine-learning pipelines
  - Avoid bias based on race, gender, age, ...
  - Reasons: laws, regulations, values, reputation, business, ...
- In scope: algorithmic fairness metrics and mitigators
  - Tabular data with protected attributes
  - Binary classification
  - Noisy trade-offs
- Important but out of scope: societal issues
  - No consensus on which techniques are right
  - Conflicting world views
  - This talk lists options, but the choice is yours

# Fairness and Data

# credit-g Dataset


y		X									
label	protected attributes										
class	personal_status	age	checking_status	duration	credit_history	purpose	credit_amount	savings_status	employment	...	
0	good	male single	67.0	<0	6.0	critical/other existing credit	radio/tv	1169.0	no known savings	>=7	...
1	bad	female div/dep/mar	22.0	0<=X<200	48.0	existing paid	radio/tv	5951.0	<100	1<=X<4	...
2	good	male single	49.0	no checking	12.0	critical/other existing credit	education	2096.0	<100	4<=X<7	...
3	good	male single	45.0	<0	42.0	existing paid	furniture/equipment	7882.0	<100	4<=X<7	...
4	bad	male single	53.0	<0	24.0	delayed previously	new car	4870.0	<100	1<=X<4	...
...	...	...	...	...	...	...	...	...	...	...	...
995	good	female div/dep/mar	31.0	no checking	12.0	existing paid	furniture/equipment	1736.0	<100	4<=X<7	...
996	good	male div/sep	40.0	<0	30.0	existing paid	used car	3857.0	<100	1<=X<4	...
997	good	male single	38.0	no checking	12.0	existing paid	radio/tv	804.0	<100	>=7	...
998	bad	male single	23.0	<0	45.0	existing paid	radio/tv	1845.0	<100	1<=X<4	...
999	good	male single	27.0	0<=X<200	45.0	critical/other existing credit	used car	4576.0	100<=X<500	unemployed	...

1000 rows × 21 columns

<https://github.com/ibm/lale>

# Fairness Meta-Data


```
fairness_info = {  
    'favorable_labels': ['good'],  
    'protected_attributes': [  
        { 'feature': 'personal_status',  
          'reference_group': ['male div/sep', 'male mar/wid', 'male single']},  
        { 'feature': 'age', 'reference_group': [[26, 1000]]}]
```



# Groups and Intersections

```
fairness_info = {  
    'favorable_labels': ['good'],  
    'protected_attributes': [  
        { 'feature': 'personal_status',  
          'reference_group': ['male div/sep', 'male mar/wid', 'male single']},  
        { 'feature': 'age', 'reference_group': [[26, 1000]]}]
```

Groups based on  
binary encoding of  
protected attributes  
and outcomes



# Axiomatic Assumptions

- Source: Friedler/Scheidegger/Venkatasubramanian. *The (Im)Possibility of Fairness: Different Value Systems Require Different Mechanisms for Fair Decision Making*. CACM 2021.
- WAE (We're All Equal)
  - All groups are essentially the same
  - If groups differ in the dataset, that is caused by structural bias
  - Motivates group fairness
- WYSIWYG (What You See Is What You Get)
  - Features and labels in the dataset accurately reflect construct space
  - Motivates individual fairness
- Algorithms cannot guarantee both WYSIWYG and WAE

# Protected Attributes Influence Outcomes

	feature	importance	std
sklearn.inspection.permutation_importance	0 checking_status	0.086000	0.011189
	1 credit_amount	0.076400	0.004923
	2 duration	0.063600	0.006681
	3 credit_history	0.053600	0.006681
	4 purpose	0.044800	0.002713
	5 age	0.027600	0.002154
	6 savings_status	0.021800	0.004534
	7 other_parties	0.014800	0.001166
	8 other_payment_plans	0.014600	0.002332
	9 residence_since	0.014000	0.002608
	10 personal_status	0.009400	0.001020
	11 employment	0.009200	0.003709
	12 housing	0.007600	0.001497
	13 job	0.006400	0.001744
	14 property_magnitude	0.005600	0.002577
	15 existing_credits	0.005600	0.001356
	16 installment_commitment	0.004400	0.002332
	17 own_telephone	0.004200	0.002135
	18 foreign_worker	0.003400	0.000490
	19 num_dependents	-0.001800	0.001327



# Other Attributes can Predict Protected Attributes

- Derived dataset where the binarized protected attribute is the target and removed from features
- Balanced accuracy 0.612 to predict personal\_status group
- Balanced accuracy 0.640 to predict age group
- Redaction would avoid disparate treatment but not disparate impact

personal_status			
	feature	importance	std
sklearn.inspection.permutation_importance	0	credit_amount	0.055200 0.007250
	1	age	0.051200 0.007250
	2	employment	0.038000 0.006481
	3	purpose	0.036600 0.001200
	4	housing	0.032800 0.001720
	5	installment_commitment	0.024200 0.002926
	6	residence_since	0.018800 0.003187
	7	num_dependents	0.018400 0.004224
	8	checking_status	0.012600 0.003200
	9	duration	0.011800 0.001939
	10	credit_history	0.008400 0.002154
	11	savings_status	0.005600 0.002245
	12	existing_credits	0.004600 0.002245
	13	property_magnitude	0.004400 0.003611
	14	job	0.002600 0.002059
	15	own_telephone	0.002200 0.000400
	16	other_payment_plans	0.002000 0.001673
	17	other_parties	0.000800 0.000400
	18	foreign_worker	0.000400 0.001625

age			
	feature	importance	std
sklearn.inspection.permutation_importance	0	housing	0.052000 0.005550
	1	personal_status	0.034800 0.007679
	2	credit_amount	0.031200 0.001470
	3	employment	0.025000 0.003406
	4	purpose	0.014400 0.003611
	5	credit_history	0.012000 0.001414
	6	job	0.012000 0.003578
	7	checking_status	0.010200 0.001720
	8	residence_since	0.009200 0.002786
	9	foreign_worker	0.006600 0.001855
	10	num_dependents	0.006200 0.003059
	11	own_telephone	0.006200 0.004956
	12	duration	0.005200 0.003187
	13	property_magnitude	0.004400 0.002577
	14	installment_commitment	0.002600 0.000800
	15	existing_credits	0.002200 0.001720
	16	savings_status	0.001600 0.001356
	17	other_payment_plans	0.000800 0.000400
	18	other_parties	0.000200 0.000400

# Fairness and Metrics

# Scikit-Learn Metrics and Scoring APIs

```
# scorer
ba_scorer = sklearn.metrics.make_scorer(sklearn.metrics.balanced_accuracy_score)
ba_scorer(trained, test_X, test_y)

# cross-validation
sklearn.model_selection.cross_val_score(
    sklearn.base.clone(trainable), train_X, train_y,
    scoring=ba_scorer, cv=StratifiedKFold(3))

# grid search
grid_search = sklearn.model_selection.GridSearchCV(
    sklearn.base.clone(trainable),
    param_grid={
        "gradientboostingclassifier__n_estimators": [1, 10, 100, 1000]},
    scoring=ba_scorer, cv=StratifiedKFold(3))
grid_search = grid_search.fit(train_X, train_y)
```

# Fairness Metrics

Metric	Formula	Inputs	Ideal	Thresholds
Disparate impact	$\Pr(y = 1 \mid X_p = 0) / \Pr(y = 1 \mid X_p = 1)$	Protected attributes $X_p$ , Labels $y$	1	Unfair to $X_p = 0$ : $<0.8$ Unfair to $X_p = 1$ : $>1.25$
Symmetric disparate impact	$di$ <b>if</b> $di \leq 1$ <b>else</b> $1/di$	Protected attributes $X_p$ , Labels $y$	1	Unfair: $<0.8$
Statistical parity difference	$\Pr(y = 1 \mid X_p = 0) - \Pr(y = 1 \mid X_p = 1)$	Protected attributes $X_p$ , Labels $y$	0	Unfair to $X_p = 0$ : $<-0.1$ Unfair to $X_p = 1$ : $>+0.1$
Equal opportunity difference	$\Pr(\hat{y}=y=1 \mid X_p = 0) - \Pr(\hat{y}=y=1 \mid X_p = 1)$	Protected attributes $X_p$ , Ground-truth labels $y$ , Predicted labels $\hat{y}$	0	Unfair to $X_p = 0$ : $<-0.1$ Unfair to $X_p = 1$ : $>+0.1$
Theil index	$E(\hat{y} - y + 1)$	Ground-truth labels $y$ , Predicted labels $\hat{y}$	1	Too much benefit: $>>1$ Too little benefit: $<<1$

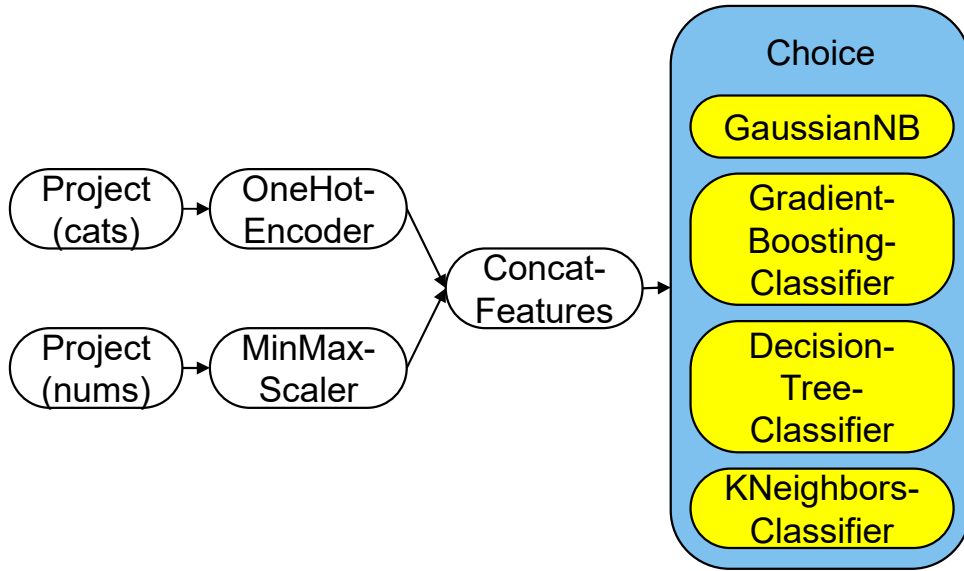
# Scikit-Learn Compatible Fairness Metrics

```
# dataset fairness meta-data
fairness_info = {
    'favorable_labels': ['good'],
    'protected_attributes': [
        { 'feature': 'personal_status',
          'reference_group': ['male div/sep', 'male mar/wid', 'male single']},
        { 'feature': 'age', 'reference_group': [[26, 1000]]}]


# scorer
di_scorer = lale.lib.aif360.symmetric_disparate_impact(**fairness_info)

# ... use like other scorer for things like cross-validation, grid-search, ...
```

# Pipelines without Mitigators



Error bars based on 10 random 75% : 25% splits



# Fairness and Pipelines

# Bias Mitigators

Mitigator	Kind	Description	Reference
Disparate impact remover	Pre-estimator	Separately shift distribution of each non-protected attribute so it is not correlated with protected attributes	Feldman/Friedler/Moelle/Scheidegger/Venkatasubramanian 2015
Reweighting	Pre-estimator	Increase sample weights for training data instances so the groups have equal total positive instance weight	Kamiran/Calders 2012
Calibrated equalized odds postprocessing	Post-estimator	Randomly flip some predictions near the decision boundaries based on group membership	Pleiss/Raghavan/Wu/Kleinberg/Weinberger 2017
Reject-option classification	Post-estimator	Deterministically flip some predictions near the decision boundaries based on group membership	Kamiran/Karim/Zhang 2012


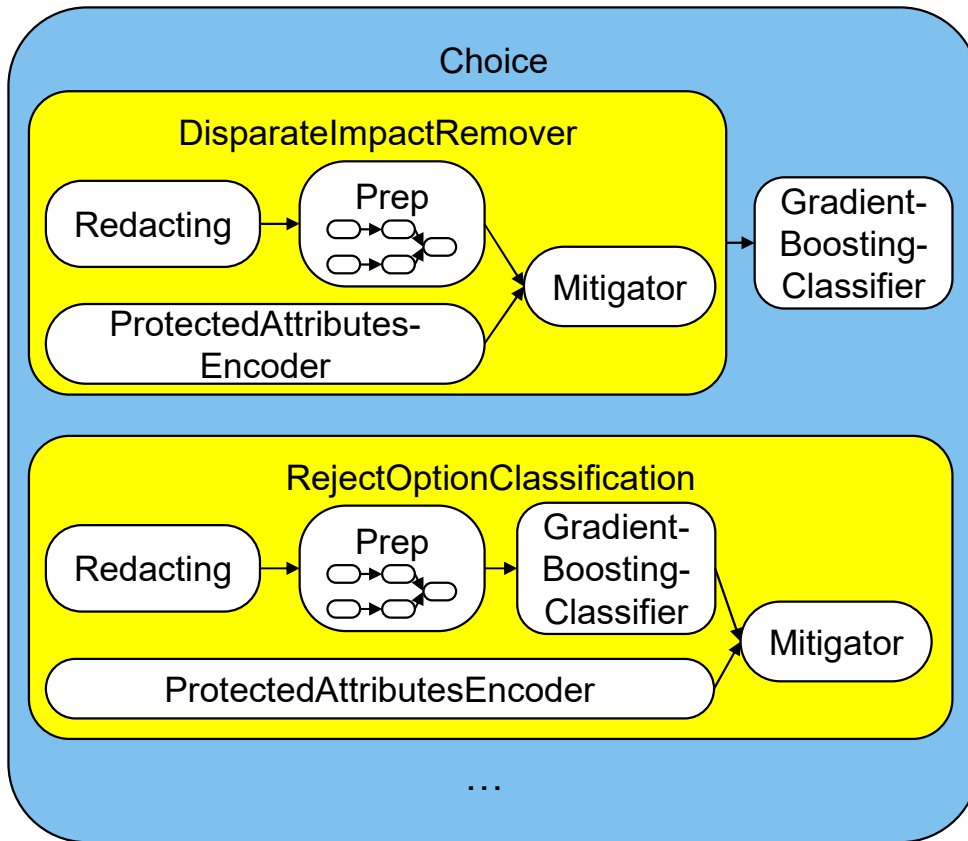


# Bias Mitigators in Python

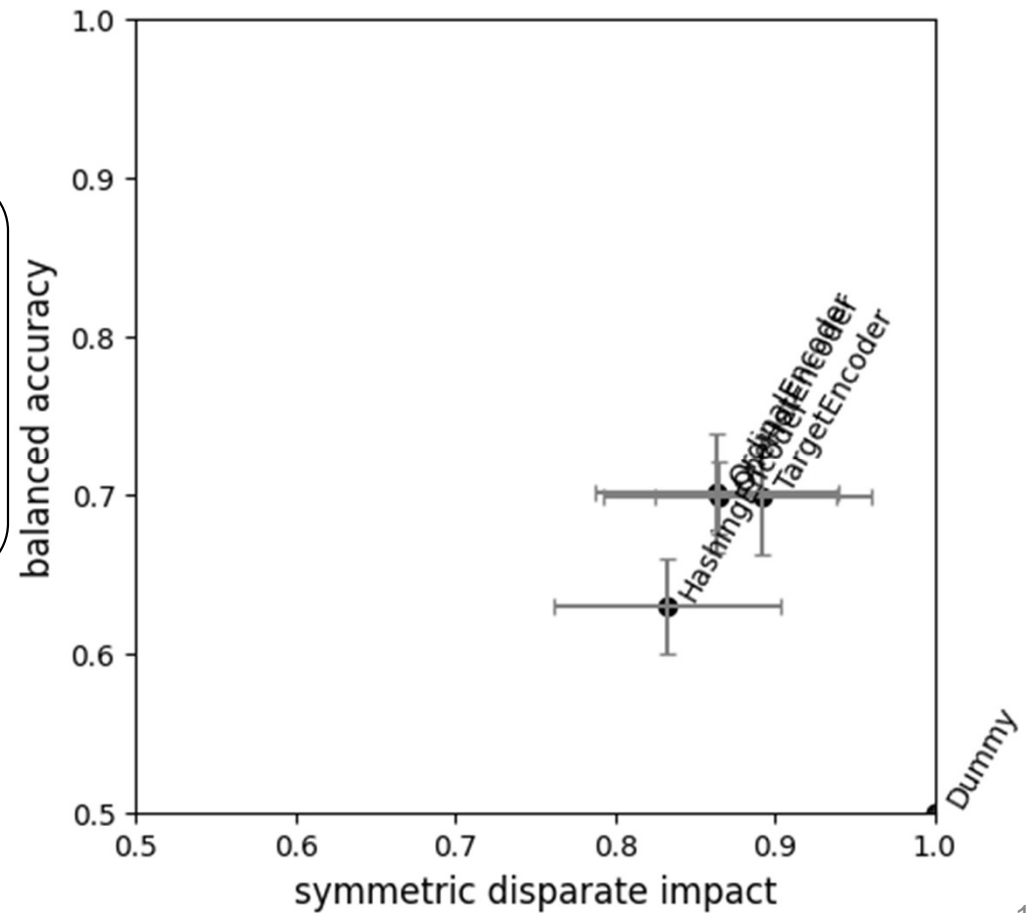
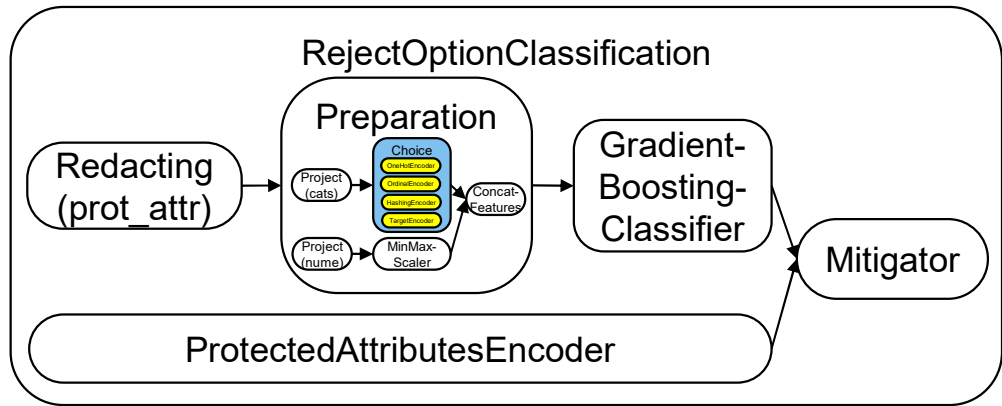
```
# dataset fairness meta-data
fairness_info = {
    'favorable_labels': ['good'],
    'protected_attributes': [
        { 'feature': 'personal_status',
          'reference_group': ['male div/sep', 'male mar/wid', 'male single']},
        { 'feature': 'age', 'reference_group': [[26, 1000]]}]

# pipeline with bias mitigator, data preparation, and final estimator
trainable = lale.lib.aif360.DisparateImpactRemover(
    **fairness_info,
    preparation=(
        (Project(columns=cat_columns) >> OneHotEncoder(handle_unknown="ignore"))
        & (Project(columns=num_columns) >> MinMaxScaler())
    ) >> ConcatFeatures
) >> GradientBoostingClassifier()
```


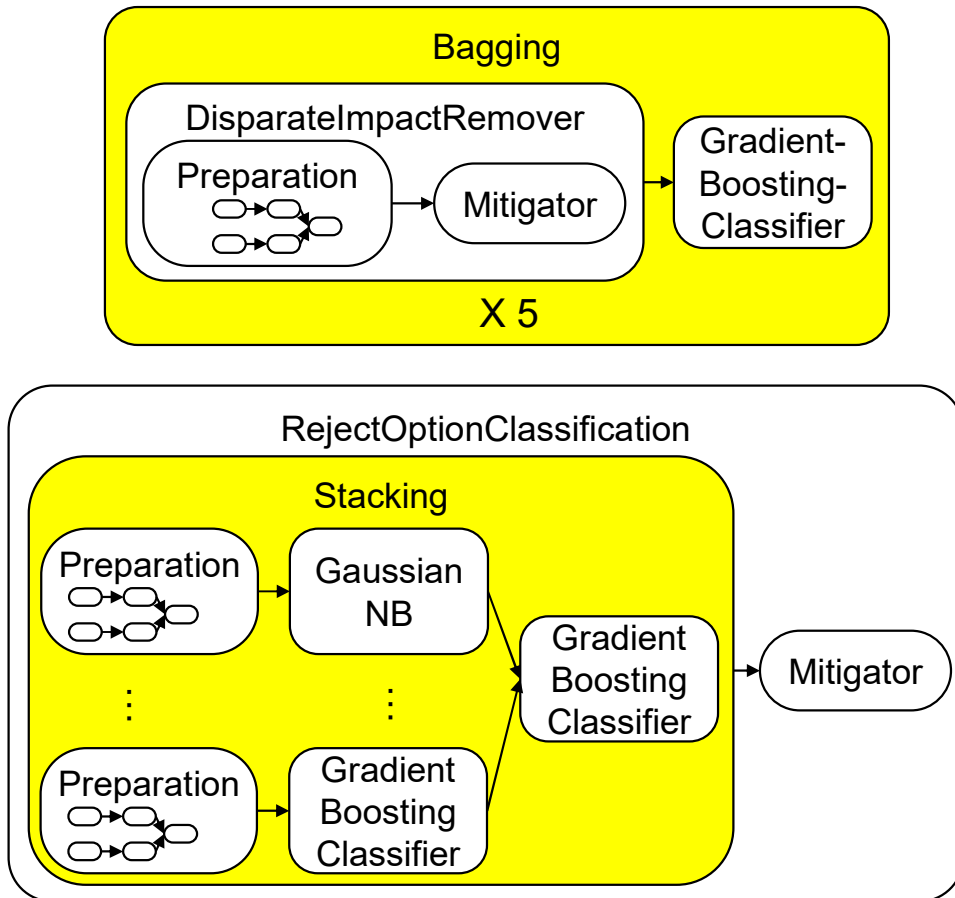
# Bias Mitigators



# Mitigators and Preprocessing



# Mitigators and Ensembles



# Fairness and AutoML

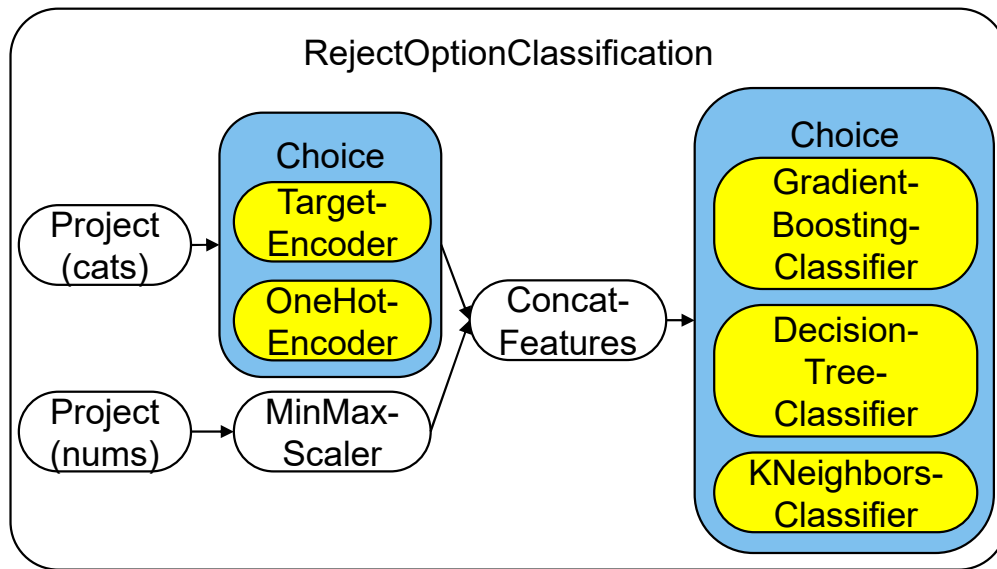
# Challenges for Fairness and AutoML

Challenge	Solution in this talk	Other solutions
Multiple objectives	<ul style="list-style-type: none"><li>• Blend into single objective via harmonic mean</li><li>• Show x-y scatter</li></ul>	<ul style="list-style-type: none"><li>• Multi-objective optimizer</li><li>• Different blending strategies</li><li>• Maximize one objective and threshold the other</li></ul>
Noise	<ul style="list-style-type: none"><li>• Show error bars</li><li>• Repeated k-fold cross validation</li><li>• Stratification by both outcomes and protected attribute groups</li></ul>	<ul style="list-style-type: none"><li>• Use larger ensembles</li></ul>

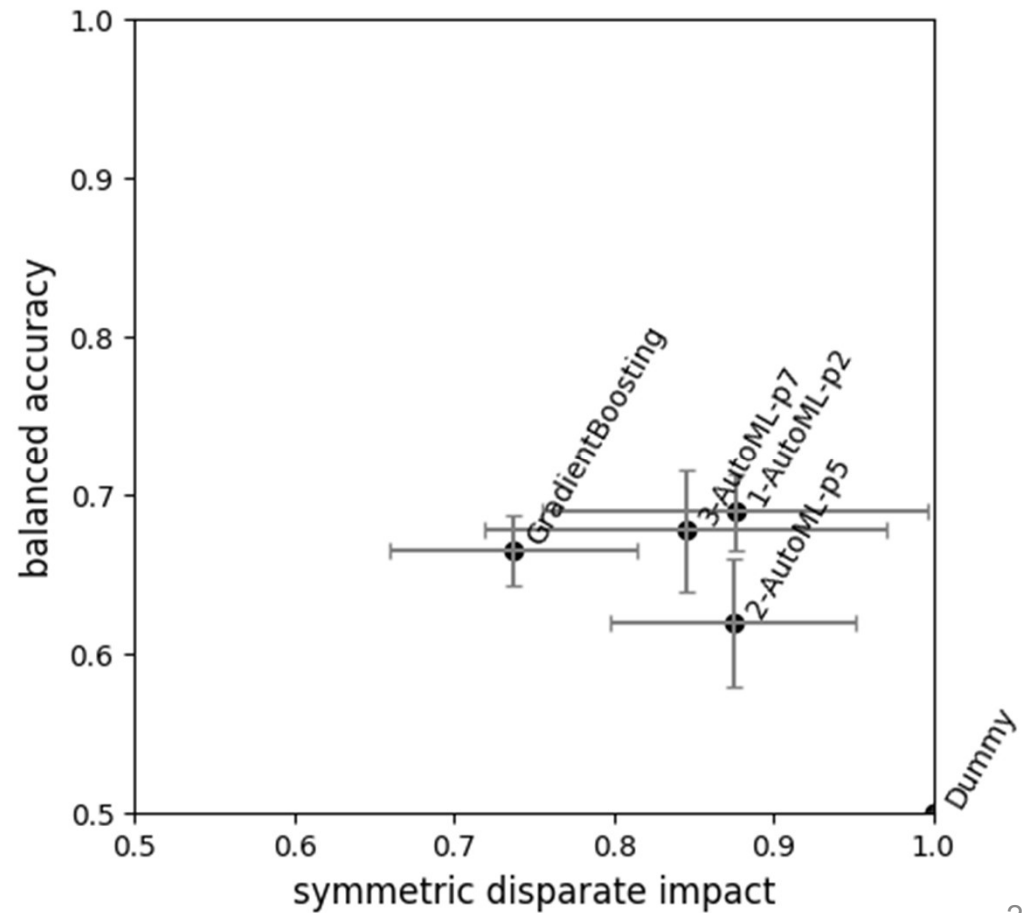
# AutoML Search Space in Python

```
planned_est_choice = lf.RejectOptionClassification(  
    **fairness_info,  
    estimator=  
        (  
            (  
                Project(columns=cat_columns)  
                >> (TargetEncoder | OneHotEncoder(handle_unknown="ignore"))  
            )  
            & (  
                Project(columns=num_columns)  
                >> MinMaxScaler  
            )  
        )  
        >> ConcatFeatures  
        >> (  
            DecisionTreeClassifier  
            | KNeighborsClassifier  
            | GradientBoostingClassifier  
        )  
    )
```

# Mitigators and AutoML



Rank	Name	Encoder	Estimator
1	p2	TargetEncoder	GradientBoostingClassifier
2	p5	TargetEncoder	KNeighborsClassifier
3	p7	OrdinalEncoder	GradientBoostingClassifier





# Conclusion

- Fairness value judgements are beyond the scope of this talk
  - E.g., WYSIWYG vs. WAS, disparate treatment vs. disparate impact, ...
  - Useful to know algorithms for accomplishing fairness goals
- This talk discussed various fairness metrics and bias mitigators
  - Fairness and data
  - Fairness metrics
  - Bias mitigators
  - Fairness and AutoML
- Lale library: <https://github.com/ibm/lale>
  - Try it out
  - Contribute